

bp

Blueprint Planner-Agent

Plan first.
Route smarter.
Ship with agents.

Terminal AI planning harness that turns vague software ideas into dependency graphs, model-routed handoffs, and validation contracts for coding agents.

github.com/marcellopps283/cli-planner-agent

```
$ blueprint
```

```
planner      gpt-5.5 / high reasoning
providers    OpenAI Codex · Claude
             Code · Gemini CLI
output       .blueprint/tasks/*.md
```

PROBLEM

AI agents execute fast. Planning still breaks.

Large AI-assisted coding projects fail when the first prompt is vague and the work is not sliced cleanly.

- Context drifts across chats and tools.
- Tasks overlap or miss dependencies.
- Expensive models get used for cheap work.
- Weak models get assigned to high-risk architecture.
- Outputs are hard to audit before code changes begin.

SOLUTION

The planning brain before worker agents start.

Blueprint runs an investigative terminal chat, validates missing requirements, routes work to exact models, and writes handoffs that other agents can execute.

Chat-first planning

Brainstorms and fills gaps before architecture is locked.

Compact context

Reads project inventory instead of dumping the entire repo.

Task graph

Creates dependency-aware work units with blocking rules.

Worker handoffs

Generates XML task prompts, model choices, and acceptance contracts.

One terminal surface from setup to handoff.

The harness keeps configuration, brainstorming, preview, and artifact generation in one flow.

1

Onboard project, providers, models, and reasoning effort.

2

Brainstorm until scope, constraints, and validation are clear.


3

Preview dependency graph and model assignments.

4

Generate architecture docs, XML tasks, and integration guide.

Planning chat, not a rigid form.



Routes by exact model, not just provider.

Risk & Fit

Architecture, security, and multi-file work require stronger reasoning.

Cost & Latency

Small documentation or formatting tasks can use cheaper, faster models.

Context & Effort

Model context window and reasoning settings are part of routing.

```
preview: task -> model -> rationale
```

```
task-001      docs polish -> gemini flash-lite
```

```
task-004      routing contract -> gpt-5.5
```

```
task-006      integration validation -> gpt-5.5 or opus
```

Versionable handoffs for worker agents

Generated handoffs

Versionable artifacts for worker agents

.blueprint/

architecture.md

dependencies_graph.json

integration_guide.md

assumptions.md

risks.md

tasks/

004-routing-contract.md

tasks/004-routing-contract.md

suggested_model: gpt-5.5

risk_level: high

allowed_paths: src/plan.ts, src/models.ts

<task_objective>

<context_rules>

<execution_prompt>

<acceptance_contract>

- Model-routed tasks
- Dependency graph
- Worker prompts

Working planner-only harness.

Version 1 plans and generates rigorous handoffs for human-controlled workers. Version 2 can become the supervisor that runs those workers from the same terminal.

- TypeScript CLI/TUI with Ink and Commander.
- Provider registry and selected model pool.
- Deterministic and LLM planning engines.
- Handoff linting, export, and smart revise.
- CI passing with 111 tests.

links

demo marcellopps283.github.io/cli-planner-agent

repo github.com/marcellopps283/cli-planner-agent